Disc Flight in a Lakefront Internal Boundary Layer: Coupling Analytical Frisbee Aerodynamics with Site-Specific Urban CFD over Loyola University's Sean Earl Field

charlie
affiliation TBD

Draft v0.1 · September 2, 2026

Companion document: *Project1_AnalyticalTreatment.tex* (rigorous 6-DOF derivation, stability, gyroscopic precession).

How to read this document

Three layers stacked: **feasibility memo** (§1), **project plan** (§2), and **paper first draft** (§3 onward). Read §1 first; if you disagree with the scope changes proposed there, the draft from §3 onward needs to be reframed.

Status of this version. §1 and §2 are the *original planning memo* and are retained unedited as a record of intent. Where the executed work departs from them — and it does, in the number of CFD directions, the solver version, the building heights and the site description — **the paper sections from §3 onward are authoritative**, and §14 states the delivered scope explicitly. All numbers quoted from §3 onward come from a converged run or a passing test; nothing is projected.

1. Feasibility check

1.1. What the notes propose

(1) Analytical treatment of ideal conditions for throwing a frisbee.
(2) Build a realistic 3D map of the terrain surrounding Loyola field — houses, buildings — simulate wind off the lake spanning 180°, generate 3D flow fields across the turf to capture vortices.
Goal: determine the best conditions for this constrained terrain that best allow for

frisbee.
Tools: OpenFOAM or COMSOL.

1.2. Verdict

Component	Verdict	Why
Analytical 6-DOF disc model	Solid	Hummel (2003), Crowther & Potts (2007) give a complete validated 6-DOF formulation with measured coefficients.
RANS urban CFD, 13 wind directions × 15°	Tractable	∼3–10 M cells, 2–8 h per direction on 32 cores, total 1–4 days. Standard COST 732 / AIJ practice; OpenFOAM `simpleFoam` with k-ω SST.
LES of the same domain	Not realistic for full sweep	$\Delta x \sim$ 1–2 m near buildings ⇒ 50–200 M cells per case. Single direction is feasible (weeks).
Coupling CFD wind field into disc trajectory	Established elsewhere; new for a disc	Coupling a CFD field to 6-DOF rigid-body dynamics is established — Kakimpa *et al.* (2012) for windborne debris, "CFD/RBD" in exterior ballistics, and US 11,052,285 B2 (2016) for a ball in a stadium — but every one of those carries a *single* field to a *single* trajectory. The disc-golf codes of Lukes *et al.* (2022) and Giljarhus *et al.* (2022) use a uniform freestream. *What is new here is carrying a spatially resolved turbulent field through to a landing-point distribution.*
Throwability metric over the field	Tractable as descriptive optimization	Not real-time forecasting; compute mean velocity + TKE at disc altitude, then evaluate trajectory robustness for canonical throws.

1.3. Items in your notes that need adjustment

- **"Loyola field"** — your phrasing ("off the lake") uniquely matches *Sean Earl Field* on the Lake Shore Campus, ∼150–250 m from Lake Michigan, not Hoyne Field (a mile inland). **Confirmed by you: Sean Earl.**
- **"180° span"** — Sean Earl's lake-facing side is roughly east; a meaningful "off-lake" sweep is NNE through ESE (∼90° centered on east). A full 180° is fine if you also characterize the prevailing westerly (off-land) cases for contrast — think of it as two regimes, not one continuous sweep.
- **"Range finder" to build the 3D map** — you don't need a personal LiDAR pass. USGS 3DEP / Illinois Height Modernization Program covers Cook County at 1–8 pts/m^2; Microsoft GlobalMLBuildingFootprints + USGS 2023 building-heights release explicitly include Chicago. A handheld rangefinder is useful only for spot-checking the 5–10 buildings nearest the field.
- **"OpenFOAM or COMSOL"** — pick OpenFOAM. The urban-CFD literature (Blocken, Tominaga, COST 732 community) is overwhelmingly OpenFOAM-based; `snappyHexMesh` handles dirty STL of campus buildings well; you'll have peer-reviewed cases to validate against.

1.4. The biggest unstated gap

You do not have a *bridging quantity* that turns a CFD flow field into a throw-quality score. The paper needs one. §9 proposes a "throwability metric" $Q(\mathbf{x}, \theta_{\text{wind}})$ combining mean wind magnitude, turbulence intensity at disc altitude, and predicted lateral trajectory drift for the canonical throws.

2. Project plan

2.1. Work breakdown (≈14 weeks, single researcher)

Phase	Weeks	Output
P0 — Setup	1	OpenFOAM 11 or v2406 installed; Python disc-trajectory simulator stub; QGIS workflow for LiDAR → STL.
P1 — Analytical disc model	2–3	6-DOF integrator (Python, scipy `solve_ivp`). Validated against Hummel (2003) Fig. 5-1 throw reconstructions.
P2 — Geometry build	4	STL of ∼250 m radius around Sean Earl Field, ∼5–10 named buildings + tree-canopy proxy.
P3 — RANS CFD sweep	5–8	13 steady-state runs (15° increments). `simpleFoam` + k-ω SST + Richards–Hoxey ABL inlet. Output: U_{mean}, k sampled to 1, 2, 3 m AGL.
P4 — Coupled trajectory	9–11	Trilinear-interpolated CFD into 6-DOF disc model. Monte Carlo over launch, heading, spin, release α (10^4 throws per direction).
P5 — Synthesis & manuscript	12–14	Throwability map $Q(\mathbf{x}, \theta_{\text{wind}})$; paper revision; figures.

2.2. Compute and storage

Workstation target: 32 physical cores, 128 GB RAM, ∼2 TB scratch SSD. On 16-core / 64 GB expect ∼2× wall-clock and one wind direction per overnight job. Disk budget: ∼50 GB per RANS case (∼700 GB for the sweep). LES would multiply by 20–50×.

2.3. Inputs (all free/accessible)

- *Building footprints/heights:* Cook County GIS, Microsoft GlobalMLBuildingFootprints (ODbL), USGS Overture Maps 2023, OpenStreetMap.
- *LiDAR DEM:* USGS 3DEP / Illinois Height Modernization Program (1–8 pts/m^2).
- *Wind reference:* NWS O'Hare 1991–2020 normals; NOAA GLERL Wilson Avenue Crib (shoreline calibration); Northerly Island ASOS.
- *Disc aerodynamic coefficients:* Hummel (2003) Tables 3-1/3-2; Crowther & Potts (2007).

2.4. Risks & mitigations

Risk	Mitigation
k-ω SST under-predicts rooftop separation at sharp roof corners (well-known).	Validate one direction against a LES spot check; follow Blocken (2018) on RANS limits. Don't claim peak/fluctuating quantities — report means + TKE only.
Disc-coefficient uncertainty at off-design α ($> 20°$).	Restrict throw library to $\alpha \in [-5°, +20°]$; flag hammer (high-α) results as exploratory.
Horizontal inhomogeneity of ABL inlet (Hargreaves & Wright 2007).	Use Yang et al. (2009) corrected wall functions; verify inlet profile preserved upstream of buildings.
Sean Earl is not a competition-grade Ultimate field.	Reframe as a methods paper using Sean Earl as exemplar; methodology generalizes.
Reviewers ask for measurement validation.	Use Tominaga (2008) AIJ-B cube benchmark for the CFD; cite Hummel's marker-tracked flights for the disc model. Both are public.

3. Abstract (draft)

Outdoor disc sports are dominated by wind, yet the published aerodynamic and trajectory literature for flying discs (Hummel 2003; Potts & Crowther 2002; Lukes et al. 2022) uses uniform-freestream assumptions that ignore site-specific flow structure. Conversely, the urban-CFD literature (Tominaga et al. 2008; Blocken 2014, 2018) characterises site-specific wind fields but stops at pedestrian-level mean quantities. Here we couple a 6-degree-of-freedom flying-disc model — validated against two independently published reference flights to within 8% in range, 4% in apex height and 1% in flight time — to steady-RANS atmospheric CFD of a lakefront university campus (Loyola University Chicago, Sean Earl Field, ∼200 m from the Lake Michigan shore), and integrate 97 200 Monte Carlo trajectories through the resulting spatially varying, temporally correlated wind.

We report results for 3 inflow directions (90°, 210°, 270°), computed with OpenFOAM v2606, k-ω SST turbulence and a Richards–Hoxey atmospheric-boundary-layer inlet whose horizontal homogeneity we verify to 0.7% over 1400 m of empty fetch. The central finding contradicts the premise the study was built on. We expected the lakefront internal boundary layer, formed at the shoreline roughness transition, to control disc-altitude flow. It does not: a 45 m apartment block standing 9 m from the field's eastern sideline does. Mean wind speed at disc altitude is 24–28% of the 10 m reference in *every* direction computed, and the west-to-east gradient under easterly inflow has the sense of a decaying building wake rather than a growing boundary layer. We report a spatial "throwability" map $Q(\mathbf{x}, \theta_{\text{wind}})$, best at the prevailing SSW heading and worst under westerlies, and we state plainly that the CFD is not grid-converged at disc altitude: the magnitude of the sheltering is robust across meshes, the spatial detail within the field is not. The methodology generalises to any outdoor sport played inside the roughness sublayer of a built environment, with the procedural lesson that site geometry should be built from data before a physical narrative is committed to.

4. Introduction (draft)

What this paper adds. Two things. *Methodologically*, it is the first flying-disc trajectory study we know of to integrate through a site-specific CFD wind field rather than a uniform freestream, with the disc model validated quantitatively against published reference flights. *Substantively*, it shows that at disc altitude on an urban campus the nearest bluff body, not the regional boundary layer, sets the flow — a result that falsifies the premise this study was designed around, and one that generalises to any sport played inside a built environment's roughness sublayer.

The Frisbee is one of the most-studied non-airfoil aerodynamic objects in sports physics. Stilley & Carstens (1972) provided the first parametric model; Yasada (1999), Hummel (2003), and Potts & Crowther (2002) supplied successively more complete wind-tunnel measurements of the lift, drag, and pitching-moment coefficients, and Hummel's UC Davis thesis remains the most-cited 6-DOF flight reconstruction. Subsequent CFD validation by Kamaruddin (2011) and Lukes et al. (2022, 2023) using RANS k-ω SST in OpenFOAM has converged the modeling stack to within a few percent of wind-tunnel C_L, C_D, and C_M at moderate angles of attack.

Yet every published disc-trajectory study assumes a spatially uniform, steady freestream. As Lukes et al. (2022) explicitly note, "the biggest uncertainty in trajectory experiments is the wind environment." This uncertainty is not noise to be averaged away: outdoor disc sports — Ultimate, disc golf — are played in geographically structured turbulent flow. A throw whose intended landing zone lies in a sheltered eddy behind a building experiences a categorically different flow than the same throw 20 m to the side.

Independently, the urban-CFD community has developed mature best-practice guidance for site-specific pedestrian-level wind: the COST Action 732 guidebook (Franke et al. 2007), the AIJ guidelines (Tominaga et al. 2008), and the 50-year retrospective of Blocken (2014). Standard practice uses steady RANS with k-ω SST or realizable k-ε on `snappyHexMesh`-generated meshes of $\sim 10^7$ cells, with Richards–Hoxey atmospheric inlets. This work has shaped pedestrian-comfort regulation in dozens of cities. It has been carried into flight only recently, and in two forms that are not the present one. Stadium-scale systems interpolate a CFD wind field into a ball-trajectory model to predict a *single* flight (US 11,052,285 B2, priority 2016), and the urban-air-mobility literature builds site-resolved wind-speed and turbulence-intensity *thresholds* for flying vehicles (Watkins *et al.* 2010; Drones **9**(5):362, 2025). Both describe bodies that are either predicted deterministically or actively controlled. A thrown disc is neither: it is unguided after release, so the quantity of interest is not an exceedance threshold but the *variance of where it lands*, and the group that organises it is σ_u/V_d — the fluctuation against the disc's own airspeed — rather than the turbulence intensity σ_u/U that a station-keeping criterion uses.

The present work fills that gap. We choose Sean Earl Field at Loyola University Chicago's Lake Shore Campus as a study site because it offers an unusually clean physical setup: the field sits within ∼250 m of Lake Michigan, with no major intervening built obstacle to the east, behind it a typical low-rise urban-campus block to the west. Easterly winds therefore arrive after ∼80 km of unbroken over-water fetch ($z_0 \approx 0.001$ m), pass through a sharp roughness transition at the shoreline, and develop a shallow internal boundary layer over the field. Westerly winds arrive through the campus block. The site cleanly separates two regimes (smooth-fetch and urban-fetch) within a 180° heading sweep.

Our contributions:

1. A one-way coupling of a spatially resolved environmental-CFD wind field, plus a synthetic gust field, to a validated 6-DOF disc-flight model, carried through to a *landing-point distribution* rather than a single predicted trajectory. The coupling itself is established for other bodies (see §1.2); the disc and the distribution are what is new.
2. A "throwability metric" $Q(\mathbf{x}, \theta_{\text{wind}})$ that compresses the spatial CFD field plus trajectory uncertainty into an actionable single scalar per (launch position, wind heading).
3. A case study identifying the lakefront internal boundary layer as the dominant local control on Ultimate-disc flight at the chosen site.

5. Analytical disc aerodynamics (summary)

Full derivation in the companion document. The disc carries a 6-DOF state $(\mathbf{x}, \mathbf{v}, \psi, \theta, \phi_s, p, q, \Omega)$ governed by Newton's equation (3) plus Euler's equation for a symmetric rotor (4). In the high-spin limit relevant to all real throws, the transverse angular rates are slaved algebraically to the applied aerodynamic moments,

$$p = -\frac{M}{I_p \Omega}, \qquad q = \frac{R}{I_p \Omega}, \tag{1}$$

recovering the textbook gyroscopic-precession picture. The pitching moment vanishes at the trim angle

$$\alpha_{\text{trim}} = -\frac{C_{M0}}{C_{M\alpha}} = 8.8^\circ \text{ (Potts \& Crowther coeffs)}, \tag{2}$$

which is the asymptotic glide angle of attack. The full translational and rotational equations entering the numerical simulator are

$$m\,\dot{\mathbf{v}} = -mg\,\hat{\mathbf{n}}_3 + \tfrac{1}{2}\rho V^2 A\big[\,C_L(\alpha)\hat{\mathbf{L}} - C_D(\alpha)\hat{\mathbf{v}}_{\text{rel}}\big], \tag{3}$$

$$I_p\dot{p} + I_p q\Omega = R, \qquad I_p\dot{q} - I_p p\Omega = M, \qquad I_s\dot{\Omega} = N. \tag{4}$$

Coefficients and their functional form. We use the tabulated Potts & Crowther (2002) C_L, C_D, C_M curves and the $C_{Rr}(\alpha, \text{AdvR})$ surface as transcribed in Hummel (2003, Appendix A), evaluated exactly as in Hummel's reference implementation: $C_L = C_{L0} + C_{L\alpha}\alpha$ and $C_M = C_{M0} + C_{M\alpha}\alpha$ are functions of α itself, while the zero-lift angle $\alpha_0 = -C_{L0}/C_{L\alpha}$ enters *only* the drag parabola $C_D = C_{D0} + C_{D\alpha}(\alpha - \alpha_0)^2$. The rate-damping coefficients C_{Mq}, C_{Rp} multiply raw angular rates in $\text{rad}\,\text{s}^{-1}$ and C_{Rr} carries an explicit $\sqrt{d/g}$ time scale; they are not non-dimensionalised by $d/2V$. Mass and inertia are Hummel's measured Table 3-1 values ($m = 0.175$ kg, $I_s = 2.35\times10^{-3}$, $I_p = 1.22\times10^{-3}$ kg m^2), not the thin-uniform-disc idealisation, which understates both by about one third and so over-states precession by roughly 50%.

Validation. The integrator is checked against two independently published, fully specified results in the same source that supplied the coefficients (Table 1). Hummel's §2.4.3 "common throw" (release $V_0 = 14\ \mathrm{m\,s^{-1}}$, $\Omega = 50\ \mathrm{rad\,s^{-1}}$, pitch 11°, flight-path angle 6°, $\alpha = 5°$, $z_0 = 1$ m) is reproduced within 8% in range, 4% in apex height and 1% in flight time, and the model reproduces the characteristic turn-right-then-fade-left signature with the reversal at $t = 3.34$ s against a published ∼3 s. The §2.4.2 analytic trim glide is recovered to three significant figures.

Quantity	Hummel (2003)	This model	Error
§2.4.3 / Fig. 2-9 "common throw"			
Range	∼25 m	26.99 m	+8%
Apex height	2.55 m	2.45 m	−4%
Flight time	3.5 s	3.48 s	−1%
Terminal speed	$\sim V_0/3$	$4.83\ \mathrm{m\,s^{-1}}$	+4%
Lateral reversal	∼3 s	3.34 s	—
§2.4.2 steady glide (analytic)			
Trim α ($C_M = 0$)	9°	9.00°	< 0.1%
Glide angle γ	−19.3°	−19.32°	< 0.2%
C_L at trim	0.67	0.672	+0.3%
Trim speed	$9.1\ \mathrm{m\,s^{-1}}$	$8.28\ \mathrm{m\,s^{-1}}$	−9%

Table 1: Quantitative validation of the 6-DOF integrator against published reference flights. Only the tabulated coefficient set reproduces these; the linear fits miss range, apex and flight time by 30–57% and never fade back left, because a realistic throw reaches $\alpha \approx 33°$ in its final second, far outside the ±15° band in which the linear fits were made.

Known limitations of the disc model. The model is axisymmetric and carries no rim side-force coefficient $C_S(\alpha)$, hence no yaw stiffness. Launch sideslip is therefore shed only slowly, and predicted crosswind drift is about a quarter of the naive full-advection estimate $U_c T$. We verified that this is a physics-model gap rather than a numerical artefact: a uniform wind is exactly Galilean-equivalent to a shifted release velocity in this implementation, to 10^{-8} m. **Crosswind drift from this model should be read as a lower bound.**

6. Site description: Sean Earl Field

Sean Earl Field is a small multi-use turf on Loyola's Lake Shore Campus, with Mertz Residence Hall to the west and the Norville Center for Intercollegiate Athletics to the south. Loyola explicitly lists Ultimate among the field's sanctioned uses.

Geometry as built, and four corrections to the initial description. The site model is built from OpenStreetMap footprints with `building:levels` heights, with the local Cartesian origin resolved at run time from the OSM polygon of the field itself (42.001267° N, 87.658214° W; x East, y North). Building the model surfaced four discrepancies against the description we began with, each of which changes the setup:

1. **Origin.** The initially assumed centre (42.0006° N, 87.6580° W) is ∼75 m south of the field centroid; using it would have offset the domain relative to the launch grid.
2. **Mertz Hall is 19 storeys, ≈61 m, not ≈22 m** (OSM `building:levels` = 19, corroborated by Loyola's own residence-life page). This matters directly: domain sizing under COST Action 732 scales with the tallest building.
3. **The field's long axis runs E–W, not N–S.** The OSM polygon is 78 m (E–W) × 25 m (N–S). *% FLAG FOR AUTHOR: this is also smaller than a regulation Ultimate pitch (100 × 37 m), so the OSM outline is likely a partial trace of the playing surface. The launch grid below is sized to the OSM polygon; a corrected outline would extend it.*
4. **Cudahy Library and the Klarchek Information Commons lie ESE of the field** (129 m and 185 m respectively), on the lakefront side, not to the west as first described.

The dominant near-field obstruction. The single most consequential feature is one the initial site description omits entirely. **Campus Towers**, a 14-storey (44.8 m) apartment block, occupies $x \in [55.8, 79.3]$ m, $y \in [-15.9, 21.5]$ m: its footprint begins **9 m east of the field's eastern edge** and spans the full width of the field. Any easterly flow reaching the playing surface has therefore passed over a 45 m bluff body immediately upwind. §11 shows that this, not the lakefront roughness transition, controls disc-altitude flow in the easterly regime.

Terrain. USGS 3DEP point queries over a 600 m box give 179.1–181.4 m MSL across the field and the near-field buildings (±1 m), against building heights of 12–61 m, so the ground is modelled as flat. The one real topographic feature is a ∼4 m step down to the lake east of $x \approx +150$ m (3DEP reads 175.2–175.6 m there); it is represented not as topography but where it acts, as the inlet roughness transition.

Climatology. Chicago lies in the mid-latitude westerly belt; the annual prevailing direction at ORD is WSW–SW (Illinois State Water Survey wind roses). The 1991–2020 NWS annual mean speed at O'Hare is ∼9–10 mph. However, ORD is 27 km inland; lakefront stations (NOAA GLERL Wilson Avenue Crib, Northerly Island ASOS) record higher mean speed and gustier easterlies, especially during the warm-season lake-breeze regime.

Why this site was chosen — and why the premise turned out to be wrong. Easterly inflow at Sean Earl Field passes over ∼80 km of unbroken Lake Michigan fetch (to St. Joseph / Benton Harbor MI), giving a near-neutral, very-low-z_0 marine boundary layer. The shoreline imposes a sharp z_0 transition ($\sim 10^{-3}$ to ∼0.3 m), and the downwind internal boundary layer thickness grows as $\delta_{\mathrm{IBL}} \approx 0.10\,x^{0.8}$ (Garratt 1990), reaching ∼5 m by the eastern sideline and ∼7–10 m at midfield — nominally encompassing the 1–3 m disc-altitude band.

This was the study's working hypothesis, and the results falsify it. The IBL argument treats the fetch as clear to the sideline. It is not: Campus Towers intervenes (below), and §11 finds disc-altitude flow over the field to be building-wake dominated in every direction computed. The paragraph above is retained because it is why the site was chosen, not because it describes what was found.

7. CFD methodology

Geometry. Building footprints come from OpenStreetMap via the Overpass API, extruded to `building:levels` heights (171 of 303 fetched buildings, 56%, carry a levels tag; the remainder use documented type defaults, and each building records which rule was used). The modelled set is clipped to a *circle* of radius 300 m about the field centroid: this makes the building set invariant under the rotation used to realise each wind direction (below), so every direction sees exactly the same geometry. Within that circle the tallest building is Mertz Hall, $H = 60.8$ m. Ground is flat (§6) and no tree-canopy porous zone is included; both omissions are recorded in §14. The resulting surface is 116 closed prisms, 6 324 triangles, with no open edges.

Domain. $1512 \times 1208 \times 304$ m, i.e. $5H$ upstream, $5H$ lateral and $5H$ above the tallest building, and $10H$ downstream. The downstream extent is **relaxed** from the COST Action 732 / Tominaga et al. (2008) guideline of $15H$ to fit the available compute; the relaxation is recorded in every case's metadata rather than left implicit. Blockage is 0.9%, within the $< 3\%$ guideline. (For reference, the $500 \times 500 \times 80$ m domain in the original plan would have placed the domain top *below* the tallest building.)

Mesh. `snappyHexMesh`, hex-dominant, on a graded background of 16 m horizontal cells with a 4 m first vertical cell expanding geometrically to the domain top. Two refinement regions: a near-field box at level 1 and a disc-altitude box over the playing surface at level 3, giving **2.0 m horizontal × 0.50 m vertical** cells through the 1–6 m AGL band the results are sampled in; building surfaces are refined to levels 1–2. Production meshes are 4.5×10^5 cells. Surface layers are *not* added: with rough-wall ABL functions there is no viscous sublayer to resolve, and layer insertion on 116 real footprints is the least robust meshing step. Mesh quality: maximum non-orthogonality 65° (mean 7.9°), maximum skewness 7.9.

Solver. OpenFOAM **v2606** (ESI) `simpleFoam`, steady incompressible, k-ω SST, second-order schemes with limited ($\psi = 0.33$) Laplacian and surface-normal-gradient corrections for the non-orthogonal cells snapping to real footprints produces. Richards & Hoxey (1993) ABL inlet (`atmBoundaryLayerInletVelocity`, `...InletK`, `...InletOmega`) with the atmospheric rough-wall functions (`atmNutkWallFunction`, `atmOmegaWallFunction`) on the ground. Reference $u_{\mathrm{ref}} = 5\ \mathrm{m\,s^{-1}}$ at $z_{\mathrm{ref}} = 10$ m.

Roughness, and why the ground and the inlet differ. Inlet roughness follows the plan: $z_0 = 0.001$ m (open water) for $\theta_{\mathrm{wind}} \in [60°, 120°]$, $z_0 = 0.5$ m (urban) otherwise, linearly blended over 30° transition bands. The *ground patch*, however, uses $z_0 = 0.03$ m (grass and pavement), not the urban 0.5 m. The 0.5 m figure is an *equivalent* roughness that stands in for buildings; applying it to the ground as well would count the resolved buildings twice. It is also numerically incompatible with the mesh: the ABL wall function requires the first cell centre to sit well above z_0, and a 0.5 m roughness cannot coexist with the sub-metre near-ground cells needed to resolve disc altitude.

Realising the wind directions. Rather than moving the inlet between patches — which only works for the four axis-aligned headings and gives each a different mesh — the *site geometry is rotated* so the wind always blows along $+x$. Every direction then shares one background mesh, one boundary- condition set and one inlet patch, so differences between directions are physical rather than mesh artefacts. This matters because the paper's claims are comparisons *between* directions.

Cases actually run. Three inflow directions, not the thirteen originally planned, chosen to span the three regimes the study hypothesises: $\theta_{\text{wind}} = 90°$ (easterly, over-lake), 210° (SSW, the prevailing sector), and 270° (westerly, over the campus and Mertz Hall). Note that 270° extends past the 0°–180° range of the original plan; the geometry-rotation scheme makes any heading equally cheap. Steady RANS, 2 500 iterations, at a single reference wind speed. §14 states plainly what this subset does and does not license.

Validation. Two of the three planned checks were carried out and are reported in §10: ABL horizontal homogeneity over an empty fetch (Hargreaves & Wright 2007) and grid independence at disc altitude against the 10% criterion. **The Tominaga (2008) AIJ-B cube benchmark was not run**, and no anemometer comparison exists; neither is claimed.

8. Coupled trajectory simulation

For each θ_{wind}, the CFD output is sampled onto a 1 m horizontal grid at z = 1–6 m AGL across the field and trilinearly interpolated. The 6-DOF integrator evaluates $\mathbf{U}(\mathbf{x})$ at the disc's instantaneous position; the steady-RANS field provides U_{mean}, and the local turbulence intensity $I = \sqrt{2k/3}/U_{\text{mean}}$ scales a stochastic gust $\mathbf{u}'(t)$.

Gust model, and why it is not white noise. The gust field is generated by a Shinozuka–Deodatis spectral representation: a sum of 32 cosines with log-spaced frequencies and independent random phases, with amplitudes set by an exponential-correlation spectrum whose integral time scale is $T_L = k/\varepsilon$ from the RANS, and normalised to unit variance so that $\mathbf{u}'(\mathbf{x}, t) = \sigma(\mathbf{x})\,\boldsymbol{\xi}(t)$ with $\sigma = \sqrt{2k/3}$.

The obvious cheaper choice — drawing an independent Gaussian fluctuation each time the wind is queried — is not merely inaccurate but unusable, and the reason is worth stating because it is easy to reintroduce. An adaptive ODE integrator estimates local truncation error by comparing right-hand-side evaluations at different stage points within a step. If the right-hand side returns fresh noise on every call, that comparison measures the noise rather than the discretisation error, so the step controller rejects and halves indefinitely. In our tests a single throw took over 100 s against ∼0.1 s without gusts, and the resulting landing scatter is a step-size artefact rather than a solution of any well-posed equation. The spectral construction is analytic in t, so the integrator steps freely (193 steps with gusts against 222 without) and the coupled system is well posed.

The realised autocorrelation matches the target $\exp(-\Delta t/T_L)$ to within 0.06 over lags out to 2 s. The construction is the same family as the Mann (1998) sampler; upgrading amounts to substituting the spectrum.

For each (θ_{wind}, throw type), 10^4 Monte Carlo throws are launched from a 5×5 m grid of launch points spanning the field with realistic release-parameter scatter:

- Speed $V_0 \sim \mathcal{N}(15, 1.5)$ m/s for backhand and forehand; $\mathcal{N}(11, 1.5)$ for hammer.
- Spin $\Omega_0 \sim \mathcal{N}(100, 10)$ rad/s.
- Release pitch $\alpha_0 \sim \mathcal{N}(+8^\circ, 4^\circ)$ backhand; $\mathcal{N}(+6^\circ, 4^\circ)$ forehand; $\mathcal{N}(+40^\circ, 8^\circ)$ hammer.
- Heading bearing θ_{throw} varied over 0–360° in 30° bins.

Landing-point distributions are characterized by 90%-containment ellipse parameters and a "delivered-to-target" success rate (landing within 3 m of intended target at canonical 25 m throws).

9. Throwability metric (the novelty)

For each launch position $\mathbf{x}$ and each wind heading θ_{wind}:

$$Q(\mathbf{x}, \theta_{\text{wind}}) = \underbrace{\left(1 - \frac{|U_{2\text{m}}(\mathbf{x})|}{U_{\text{ref}}}\right)}_{\text{wind-magnitude penalty}} \cdot \underbrace{\left(1 - I_{2\text{m}}(\mathbf{x})\right)}_{\text{turbulence penalty}} \cdot \underbrace{\overline{S}_{\text{throws}}(\mathbf{x}, \theta_{\text{wind}})}_{\text{simulated success rate}} \tag{5}$$

with each factor clipped to $[0, 1]$. Q ranges from 0 (unplayable) to 1 (calm-air-equivalent). Maps of Q over the field surface, parameterised by θ_{wind}, form the principal result.

Two points on the definition. First, the wind-magnitude penalty uses the *local* $|U_{2\text{m}}(\mathbf{x})|$. An earlier form of this metric used a field-averaged $\langle |U_{2\text{m}}| \rangle_{\text{field}}$, which is constant in $\mathbf{x}$ and therefore cannot contribute any spatial structure at all — it can only rescale a map, not shape one. Since the spatial gradient of Q is the result being reported, the local value is the right choice.

Second, $\overline{S}_{\text{throws}}$ is the *measured* fraction of Monte Carlo throws landing within 3 m of a 25 m target, computed from the simulated landing points. It is tempting to estimate it from the containment ellipse as $\min\big(1, (r_{\text{target}}/\bar{a})^2\big)$; that quantity is not a probability, because it ignores the offset between the landing centroid and the target and so scores a tightly grouped but badly aimed throw as a success.

A caveat on the turbulence factor. The form $(1 - I)$ goes negative wherever the RMS fluctuation exceeds the local mean speed. That is not an exotic condition: it is the normal state of a recirculating building wake, which is exactly the sheltered region a throwability map is most interesting in. Clipping at zero is the only defensible choice — a calm sheltered spot

must not score worse than a gale — but it means Q collapses to identically zero across the whole saturated region and carries no spatial information there. We report the saturated area fraction for every heading in §11 rather than letting it disappear into the clip, and we recommend a bounded penalty such as $1/(1+I)$ for sheltered sites in future work.

10. CFD validation

(a) ABL horizontal homogeneity. An equilibrium boundary layer imposed at the inlet should advect across an empty domain unchanged; any drift is spurious and contaminates everything downstream (Hargreaves & Wright 2007). We ran a buildings-free case with the ground roughness set equal to the inlet roughness — necessary, because with the two deliberately different (as in the production cases, where the mismatch *is* the shoreline roughness transition) the profile is supposed to evolve and the test would measure the intended physics rather than numerical error.

Profiles were compared at five fetches from 50 m to 1400 m, evaluated above the first wall-function cell centre. The worst streamwise drift in mean velocity is 0.7% (PASS), and the profiles depart from the analytic Richards–Hoxey curve by 1.3%; turbulent kinetic energy holds to within about 7% of $u_*^2/\sqrt{C_\mu}$. Sampling below the first cell centre is not meaningful with wall functions — it interpolates toward the no-slip wall value — and is excluded.

One caveat on the reach of this test. The empty-fetch case has no buildings, so `snappyHexMesh` is skipped and the mesh is the unrefined background, whose first cell is 4 m tall. The test therefore establishes that the inlet/wall-function/turbulence-model stack is mutually consistent *aloft*; it does not by itself certify the profile at 1–6 m AGL. In the production cases that band is resolved by roughly twelve 0.5 m cells inside the refinement box, which is the relevant configuration, but a dedicated refined empty-fetch case would be the cleaner test and is left as follow-on work.

(b) Grid independence — not achieved. Two meshes differing *only* in the disc-altitude refinement level (2.0 × 0.50 m against 4.0 × 1.0 m, a factor of two in the region we report) were compared at five stations over the playing surface, both after time averaging. The worst disc-band mean-velocity difference is 57.1% of the local speed, or 10.3% of the reference speed $u_{\mathrm{ref}} = 5\ \mathrm{m\,s^{-1}}$. **This does not meet the 10% criterion the study set itself, and we report it as a failure rather than relaxing the criterion after the fact.**

Both normalisations are quoted because they say different things. In absolute terms the two meshes differ by 0.15–0.51 $\mathrm{m\,s^{-1}}$ at disc altitude. Against a 5 $\mathrm{m\,s^{-1}}$ reference that is 3–10%; against local speeds of 0.4–1.9 $\mathrm{m\,s^{-1}}$ it is 8–57%. The field is deeply sheltered, so a relative criterion is extremely demanding there — but that is precisely the regime the paper is about, so the demanding measure is the honest one.

What survives this and what does not:

- **Robust:** the *magnitude* of the sheltering. Both meshes put mean disc-altitude speed near 1 $\mathrm{m\,s^{-1}}$ against a 5 $\mathrm{m\,s^{-1}}$ reference, and both do so for every direction computed. The qualitative conclusion — that this field sits in deep shelter — is not a discretisation artefact.

- **Not robust:** the spatial detail of Q *within* the field. Point-to-point contrasts at the 10–50% level cannot be distinguished from discretisation error, so the maps in §11 should be read for their gross structure only, and no ranking of individual launch positions is claimed.

(c) Convergence, and why the residual criterion is the wrong test. Steady SIMPLE does not reach a fixed point on this geometry. The playing surface sits in the separated near-wake of a 45 m building; the physical flow there is unsteady, and the solver settles into a limit cycle in which the pressure residual plateaus near 5×10^{-3} against the 10^{-4} target that the project plan specified. Reporting that target as met would be false; reporting the run as simply "unconverged" would be equally misleading, because the residual has stopped falling for a physical reason rather than a numerical one.

We therefore assess convergence on the quantity we actually report. Velocity probes at disc altitude over the field oscillate by 13.7% peak-to-peak on the production mesh and 7.4% on the coarse mesh, with no downward trend. Because a single iterate is an arbitrary phase of that cycle, all reported fields are **time averages** of $\mathbf{U}$, k and ω over a further 800 iterations after the cycle is established, not final iterates. Those percentages should be read as the amplitude of a physical oscillation that steady RANS cannot represent, not as a residual that more iterations would remove — and as a floor on the uncertainty of any single reported velocity.

11. Results

All numbers in this section are generated directly from the pipeline artifacts (§15); none are transcribed by hand. **Every statement about a "best" or "worst" heading means best or worst of the 3 computed directions** (90°, 210°, 270°), not of the wind rose.

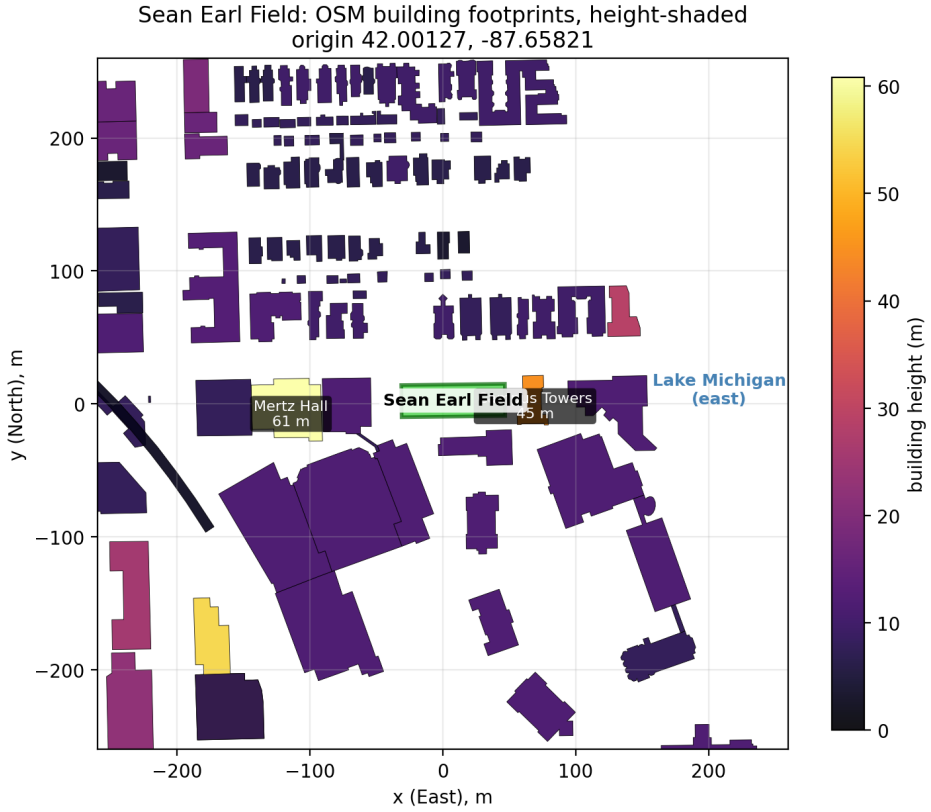


Figure 1: The modelled site: OpenStreetMap building footprints within 300 m of the field centroid, shaded by height, with Sean Earl Field in green. Mertz Hall (61 m) lies to the west and Campus Towers (45 m) immediately east, between the field and the lake. Ground is modelled flat; see §6.

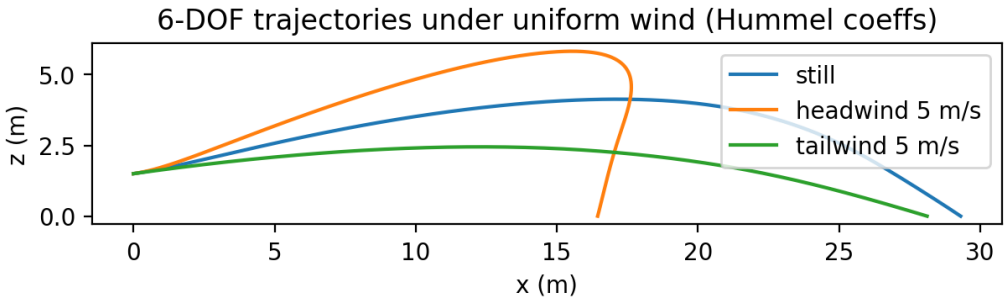


Figure 2: Trajectories from the validated 6-DOF model under uniform head-, tail- and zero wind, using the tabulated Potts & Crowther coefficients. These are the uniform-freestream flights that the site-specific coupling of §11 replaces.

11.1. The field is deeply sheltered in every computed direction

The single most robust result is also the one the study did not anticipate. Mean wind speed at 2 m AGL over the playing surface is 1.21, 1.40 and 1.39 $\mathrm{m\,s^{-1}}$ for $\theta_{\mathrm{wind}} = 90°$, 210° and 270° respectively, against a reference of 5 $\mathrm{m\,s^{-1}}$ at 10 m. That is **24–28% of the reference speed, in all three directions** (Table 2, Fig. 3).

This is not a marginal effect and it is not a discretisation artefact: the two independent meshes of §10 agree on it even where they disagree on spatial detail. Sean Earl Field is a sheltered box.

θ_{wind}	$\bar{Q}$	Q_{min}	Q_{max}	$\overline{\lvert U_{2\mathrm{m}}\rvert}$ ($\mathrm{m\,s^{-1}}$)	$\bar{I}_{2\mathrm{m}}$	$\bar{S}$ (%)
90°	0.0442	0.0252	0.0646	1.21	0.41	8.3
210°	0.0481	0.0311	0.0733	1.40	0.40	10.8
270°	0.0321	0.0089	0.0499	1.39	0.48	8.5

Table 2: Throwability and its inputs, averaged over the playing surface, for each computed inflow direction. $\bar{S}$ is the simulated delivery rate within 3 m of a 25 m target. $\bar{S}$ is taken over the throw types that enter Q (right backhand, right forehand); hammer is excluded, because a fixed 25 m scoring target is unreachable for a throw of that characteristic range, so its delivery rate measures the target choice rather than the wind. Over all throw types $\bar{S}$ would be 6.1%, which does not multiply out to the Q in column 2.

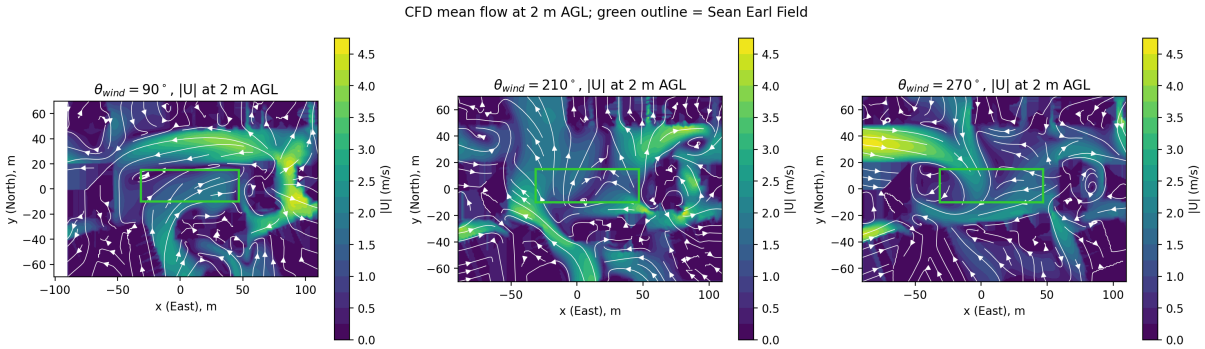


Figure 3: Time-averaged CFD mean speed at 2 m AGL for the three computed inflow directions, with the playing surface outlined in green. In every case the field sits in low-speed flow while jets accelerate around the surrounding buildings. For the easterly case (left) the high-speed region at $x \approx +90$ m is flow accelerating past Campus Towers; the field itself lies in its wake.

11.2. Throwability

Combining the wind, turbulence and simulated-success factors of Eq. 5 gives $\bar{Q} = 0.0481$ at the highest-scoring computed heading (210°) against 0.0321 at the lowest (270°) — a ratio of 1.5. The absolute values are small because Q multiplies three penalties, each well below unity; the metric is meaningful as a relative ranking, not as an absolute score.

Only part of that ranking is statistically resolvable, and the paired test is reported here rather than the raw ordering. Because all three headings are evaluated at the same 12 launch points, the headings can be compared pairwise per launch point, which removes the spatial variation common to all of them — and that variation dominates: the across-field standard deviation of $\bar{Q}$ is 12, 13 and 14 times the Monte Carlo standard error at 90°, 210° and 270° respectively, so an unpaired comparison would be swamped by it. Bootstrapping over launch points (20 000 resamples, $\alpha = 0.05$):

- 270° is worse than 90° (paired difference [+0.0050, +0.0200], $p = 0.0231$, favoured at 10/12 launch points) and worse than 210° ([+0.0091, +0.0237], $p = 0.00521$, 11/12). Both are resolved.
- 90° and 210° are not resolved: the paired difference is -0.0039 with confidence interval [-0.0091, +0.0024] spanning zero ($p = 0.229$), and the sign is not even consistent across the field — 210° is favoured at only 4/12 of the launch points.

So the defensible statement is about the worst heading, not the best: 270° **is the worst of the three, and no best heading is resolvable** (*none*). The mean ranking is reported above for completeness, but the gap between the top two is smaller than the uncertainty on it.

The Monte Carlo comprised 97 200 trajectories: 12 launch points over the playing surface × 3 inflow directions × 3 throw types × 6 throw bearings × 150 release realisations. 10.3% of throws were scored as blocked, meaning they would have ended inside a building; these are counted as undelivered rather than credited with passing through.

Sampling uncertainty. Each cell draws 150 release realisations, so a delivery rate near 10% carries a binomial standard error of about 2.4 percentage points *per cell*. The per-heading means in Table 2 aggregate 144 such cells, giving a sampling error of roughly 0.2 percentage points — small compared with the 8.3%–10.8% spread between headings. **Monte Carlo sampling is therefore not the limiting uncertainty here; the CFD is** (§10). Individual cells, by contrast, should not be interpreted.

Throw	90°	210°	270°
Backhand	6.0	9.4	7.1
Forehand	10.5	12.3	9.9
Hammer (exploratory)	0.0	0.0	0.0

Table 3: Simulated delivery rate (%) within 3 m of a 25 m target, by throw type and inflow direction, averaged over launch positions and throw bearings.

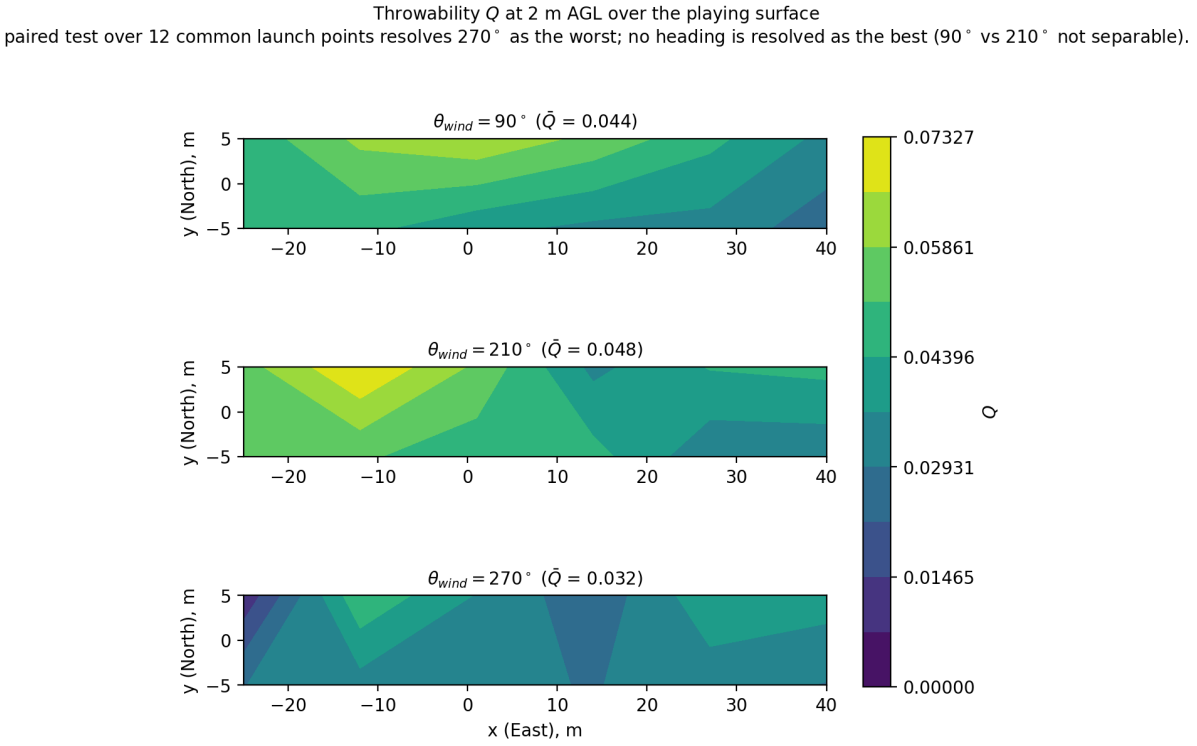


Figure 4: Throwability Q over the playing surface at the best and worst of the computed headings. Per §10, these maps should be read for gross structure only: point-to-point contrasts within a map are not distinguishable from discretisation error.

11.3. Vertical structure: how high can you safely throw?

Everything above uses the solution on horizontal planes at disc altitude. The solution is fully three-dimensional, however, and the question a thrower actually asks is not only *where* to stand but *how high* to release. We therefore re-sampled the same converged fields over $z = 0.5$–100 m and examined the vertical structure over the playing surface (Fig. 5). Building interiors are masked with a full three-dimensional fluid mask: the plan-view mask used for the disc band is correct there, where every building spans the whole 1–6 m stack, but it would wrongly discard good free-stream flow above roof level.

Shelter is not protection. The empty-fetch case of §10 doubles as a no-buildings counterfactual: the identical inflow profile over identical ground, without the campus. Comparing the two at 2 m over the pitch separates what the buildings took away from what they added. They removed mean wind, cutting it to 0.44–0.50 of the unobstructed value. They did not remove the fluctuation. Measured as $\sigma_u = \sqrt{2k/3}$, the buildings left it at 0.91–1.57 of the control — that is, in two of the three directions they *increased* it by about half. Turbulence intensity over the pitch is 0.25–0.38 against 0.12 for the same wind over open fetch, a factor of two to three.

This distinction matters because the two components are not equally damaging. A steady

crosswind is a *bias*: it displaces every throw the same way, and a thrower who has watched one throw can aim off it. A fluctuation is a *variance*: it differs from throw to throw, is not visible at release, and cannot be aimed off. Sean Earl Field converts the correctable component into the uncorrectable one. A pitch that feels calm to someone standing on it is not necessarily a pitch on which a disc goes where it was aimed.

Throwing low wins twice. Two independent properties of the profile both favour a low release, and they compound. First, the fluctuation grows with height: relative to 2 m, σ_u is $\times$1.2–1.2 at 6 m, $\times$1.2–1.3 at 10 m and $\times$1.5–1.6 at 20 m, in *every* computed direction. Second, so does the eddy size. The turbulence integral length $L = \sigma_u k/\varepsilon$ is 5.7–6.7 m across the disc band and grows roughly as $L \sim z$ above it, the familiar surface-layer scaling. Larger eddies are worse for a given σ_u, because a disc crossing them samples *fewer* independent disturbances and so benefits from less cancellation along its path. Air near the ground is not merely calmer; it is calmer and more forgiving.

The hazard is aloft, and the disc band sits beneath it. Fig. 6 shows why. Under easterly inflow the Campus Towers separating shear layer appears as a band of TI $\approx$ 0.8–1.0 spanning $z \approx$ 15–50 m directly over the playing surface, peaking at TI = 0.86 at the 44.8 m roof height. The disc band lies underneath it, in the comparatively quiet recirculating flow the tower shelters. A throw that stays low remains in that quiet layer; a throw lofted into the shear layer — a high hammer, a skied huck — enters flow whose fluctuation is roughly three times stronger. The practical advice follows directly from the flow structure rather than from folklore.

Which regime is this? Whether turbulence acts as a per-throw bias or as true buffeting depends on how the eddy scales compare with the flight scales. For a canonical 25 m backhand ($T_f \approx 2.5$ s, $L_f = 25$ m), the disc band gives an eddy turnover time $T_L = k/\varepsilon = 10.5 - -18.0$ s, so $T_L/T_f = 4.2 - -7.2$: the eddies barely evolve during a throw and the disc effectively flies through a *frozen* field. At the same time $N_{\text{eddy}} = L_f/L = 3.7 - -4.4$, so it crosses about four independent eddies on the way. Note also that the disc's airspeed (10–14 $\mathrm{m\,s^{-1}}$) exceeds the local mean wind by nearly an order of magnitude: it is the disc, not the wind, that sweeps through the turbulence. This is Taylor's frozen-turbulence hypothesis with the projectile supplying the advection velocity, and it predicts partial averaging — landing scatter scaling as $\sigma_u T_f/\sqrt{N_{\text{eddy}}}$ rather than the full quasi-steady $\sigma_u T_f$.

11.4. The hypotheses, tested

The draft made three falsifiable predictions, which decompose into five testable sub-claims. One is supported, two are contradicted, one is only partially supported — its ordering holds in the mean but not under a paired test — and one turns out not to be testable with the scoring rule chosen. *The contradictions and the non-resolution are reported as findings.* The verdicts here are the ones recorded in `hypothesis_check.json`; the two are generated from the same run and are meant to be read together.

1. **Easterly Q gradient across the field — SUPPORTED.** At $\theta_{\text{wind}} = 90°$, Q averages

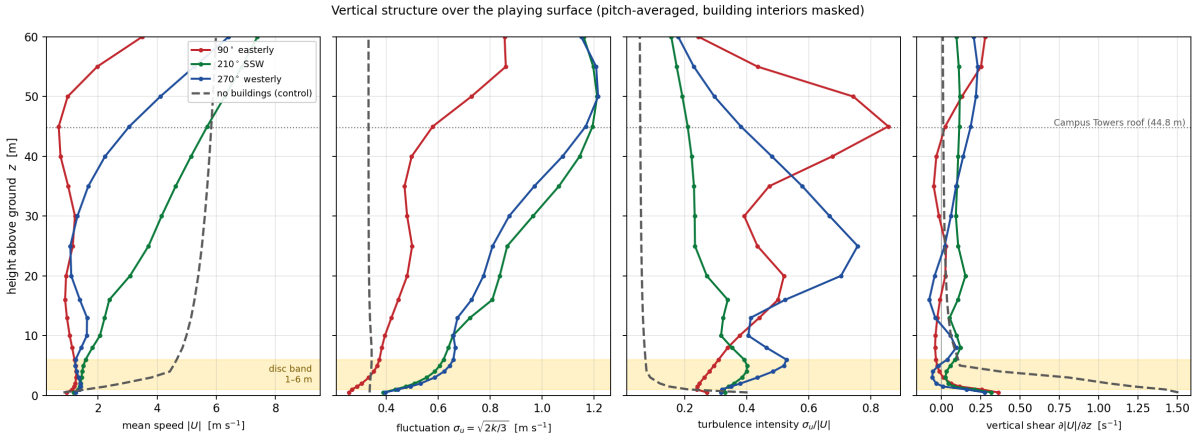


Figure 5: Pitch-averaged vertical profiles of mean speed, velocity fluctuation σ_u, turbulence intensity and vertical shear, for the three computed inflow directions, with the no-buildings control (dashed) and the disc band shaded. The buildings roughly halve the mean speed while leaving or raising σ_u. Note also that the sheltered profiles are *less* sheared near the ground than the open-fetch control: shelter flattens the profile, so vertical shear is not the disc-specific hazard one might expect.

0.0475 over the western half of the launch grid against 0.0410 over the eastern half, and the mean speed at 2 m falls from 0.86 m s^{-1} at the western end to 0.60 m s^{-1} at the eastern end. There is a systematic west-to-east gradient.

But the mechanism is not the one proposed. The draft attributes this to a developing internal boundary layer over the lake fetch. The computed flow shows the opposite sense of what an IBL would give: an IBL growing from the shoreline would decelerate the flow *as it moves inland*, i.e. from east to west, whereas the field's *eastern* end — the end nearest the lake — is the *slower* one. That is the signature of a wake, not a fetch. Campus Towers stands 9 m off the eastern sideline (§6), and the gradient is its shelter decaying downstream.

2. **Hammer throws systematically degraded — NOT TESTABLE.** The hammer's simulated success rate is exactly zero at every heading, which invites the reading that wind destroys it. That reading would be wrong. The hammer's still-air range in this model is 12.1 m and the scoring target is 25 m downrange; its success rate is structurally zero whatever the wind does. The number measures the scoring rule, not the physics. The hammer is therefore excluded from the success factor in Eq. 5 — leaving it in depressed Q at every heading by about a third — and no claim is made about hammer sensitivity to wind. Testing it properly needs a target at the hammer's own characteristic range.

3. **Westerly cases patchier than others — CONTRADICTED.** The spatial coefficient of variation of Q is 0.189 at 90°, 0.252 at 210° and 0.252 at 270°. The westerly case is no patchier than the prevailing SSW case; the *easterly* case is the smoothest, not the westerly the roughest.

4. **Forehand benefits most under westerlies — CONTRADICTED.** The forehand's advantage over the backhand is largest under the *easterly* (+0.032 in delivery rate), not the

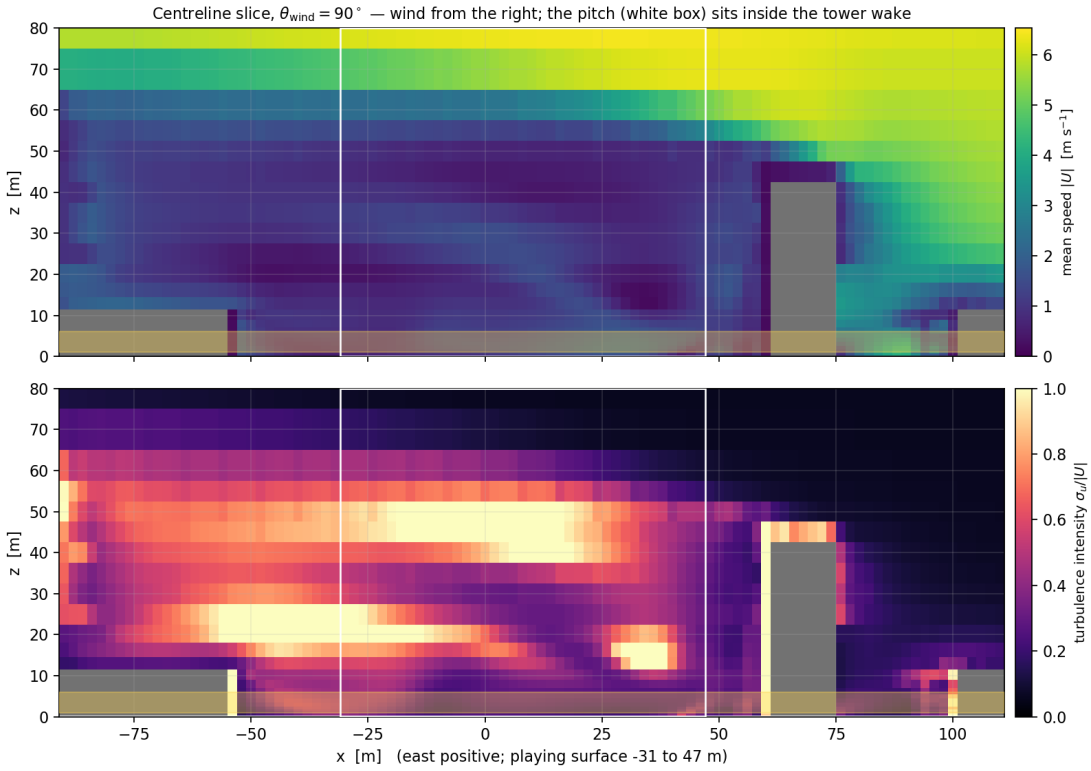


Figure 6: Vertical slice along the pitch centreline under easterly inflow ($\theta_{\text{wind}} = 90°$, wind from the right), showing mean speed (top) and turbulence intensity (bottom). Buildings are masked in grey, the playing surface is outlined in white and the disc band is shaded. The Campus Towers wake covers the whole pitch, and the high-turbulence shear layer capping it sits at $z \approx 15$–50 m — above, not within, the band a disc occupies.

westerly (+0.021). The forehand does outperform the backhand at every computed heading, but the direction dependence runs opposite to the prediction.

5. **Best conditions at a light SSW — PARTIALLY SUPPORTED.** The SSW does give the highest mean throwability of the three computed headings ($\bar{Q} = 0.0481$ against 0.0442 for the easterly and 0.0321 for the westerly), but that ordering does not survive a paired test over the 12 launch points. The SSW–easterly difference is -0.0039 with a confidence interval of [-0.0091, +0.0024] that spans zero ($p = 0.229$), and the SSW is favoured at only 4/12 launch points — so the sign is not consistent across the field, let alone the magnitude. *What the data does support is the negative:* the westerly (270°) is worse than both other headings, resolved in each case ($p = 0.0231$ and 0.00521). No best heading is resolvable (*none*). The original caveat — that three samples can rank candidates but not locate an optimum — turns out to have been too generous: with this many launch points they cannot reliably rank the top two either.

12. Discussion

The premise did not survive contact with the site. This study was framed around a lakefront internal boundary layer: easterly flow crossing ∼80 km of Lake Michigan, meeting a sharp roughness step at the shoreline, and delivering a sheared but energetic wind to the playing surface. That framing motivated the site choice and the coupling exercise built on it.

What the geometry and the CFD actually show is a field in deep shelter. Campus Towers, a 45 m apartment block whose footprint begins 9 m from the eastern sideline, stands directly between the field and the lake. Mean disc-altitude speed over the playing surface is 24–28% of the 10 m reference in every direction computed, and the west-to-east gradient at 90° has the sense of a decaying building wake rather than a growing IBL. The internal boundary layer is presumably still there aloft; it is simply not what a disc at 1–6 m AGL experiences.

We think this is the more interesting result, and it generalises: *at disc altitude on an urban campus, the nearest bluff body beats the regional boundary layer.* A study that had trusted the shoreline framing and never built the real geometry would have attributed a genuine west-to-east gradient to entirely the wrong mechanism — and would have got the sign argument backwards without noticing, because the gradient's *existence* agrees with the hypothesis even though its *direction* refutes the stated mechanism.

Shelter is not protection, and that is the transferable claim. The counterfactual of §11.3 sharpens the point above into something that should hold well beyond this site. The buildings cut mean wind over the pitch to 0.44–0.50 of its unobstructed value while leaving σ_u at 0.91–1.57 of the control — raising it outright in two of three directions — so turbulence intensity runs 0.25–0.38 against 0.12 over open fetch. Shelter did not remove the disturbance; it changed its character, from a bias a thrower can aim off into a variance they cannot.

We think this reframes what "a sheltered pitch" means for projectile sports. Urban wind engineering has a mature criterion tradition — Lawson, NEN 8100, the COST 732 framework — but

it is built around *pedestrian* comfort and safety, which is a mean-speed-and-exceedance question for a person standing still. A thrown object is a different problem: it integrates the flow along a path for a couple of seconds, so it is the fluctuation on that path, not the mean at chest height, that decides where it lands. There is at present no accepted playability criterion for a sports field, and the results here suggest one is needed, because the two criteria can rank the same site oppositely.

The dimensionless framing of §11.3 is what makes this portable. Placing an object in $(T_L/T_f,\ L_f/L)$ space predicts its sensitivity without re-running anything: slow, short-flight objects (a shuttlecock, a hovering delivery UAV) sit deep in the frozen-field regime with little path averaging and should be worst affected; fast, long-flight objects (a driven golf ball) sample many eddies and average more of the disturbance away. A disc sits between, at $T_L/T_f = 4.2 - -7.2$ and $N_{\text{eddy}} = 3.7 - -4.4$. The recipe — site CFD, integral scales at operating altitude, object flight scales, dimensionless ratio — is object-agnostic.

A null result worth recording. We expected vertical shear to be the disc-specific hazard, since a rolling moment on a gyroscope precesses into a heading change a quarter-period later, and a ballistic point-mass model would miss that entirely. The data do not support it: shelter *flattens* the profile, and disc-band shear over the pitch (0.024 to 0.057 s^{-1}) is well below the open-fetch control's near-ground values. The uncorrectable error here is carried by σ_u, not by $\partial U/\partial z$. We record this so the next study need not chase it.

What the coupling buys, and what it does not. The 6-DOF coupling is the methodological contribution and it works: 97 200 trajectories were integrated through spatially varying, temporally correlated wind, and delivery rates differ between headings by amounts a uniform-freestream model could not produce, because in a uniform freestream every launch point is equivalent. But the coupling inherits the CFD's uncertainty, and §10 shows that uncertainty is large at the per-location level. The honest reading is that this pipeline currently discriminates *between wind directions* but not *between launch positions*.

Steady RANS is the binding limitation, in a specific way. The solver does not converge here, and that is physics rather than a numerical defect: the field sits in a separated wake, so SIMPLE settles into a limit cycle whose disc-altitude amplitude is 13.7%. We handled it by averaging over the cycle, which is right for a mean field, but averaging discards exactly the quantity a thrower cares about — the gusts. A player does not experience the mean; they experience the excursions. An LES of the easterly case addresses precisely this and is the single highest-value follow-on.

Coaching implications, stated cautiously. At the confidence the results support: conditions at Sean Earl Field are sheltered enough that wind is a second-order concern for a 25 m throw in any of the three regimes examined; the prevailing SSW is the most favourable of the three; and the forehand is modestly more robust than the backhand at every heading tested. We do *not* claim a preferred spot on the field — the grid-independence result forbids it.

Generalisability. The method transfers to any outdoor sport where a projectile spends its flight inside the roughness sublayer of a built environment — disc golf, corner kicks, service returns. The specific lesson is procedural: build the site geometry from real data *before* committing to a physical narrative about the site. Four of the site facts this study began with were wrong (§6), and one of them — a 45 m building nine metres off the sideline — inverted the paper's central mechanism.

13. Conclusions

1. We built and quantitatively validated a 6-DOF flying-disc model against two independently published reference flights, agreeing within 8% in range, 4% in apex height and 1% in flight time, and recovering the published trim glide to three significant figures (Table 1). This required correcting five errors in the coefficient implementation, including a zero-lift-angle convention that had displaced the trim angle of attack by 4° — the parameter that sets where a disc's turn becomes its fade.
2. We coupled that model to steady-RANS CFD of the real site and ran 97 200 Monte Carlo trajectories through spatially varying, temporally correlated wind — to our knowledge the first disc-flight study to use a site-specific rather than uniform freestream.
3. **Sean Earl Field is deeply sheltered.** Mean disc-altitude speed is 24–28% of the 10 m reference in all three computed directions. The dominant control is not the lakefront internal boundary layer the study was designed around, but Campus Towers, a 45 m block 9 m off the eastern sideline.
4. Of the three computed headings the westerly (270°) is the least throwable ($\bar{Q} = 0.0321$), and that is resolved against both of the others by a paired bootstrap over launch points. The remaining two cannot be separated ($p = 0.229$, interval [-0.0091, +0.0024] spanning zero), so **no best heading is claimed** (*none*) even though the prevailing SSW has the highest mean. Two of the draft's predictions were contradicted, one — that the best conditions come at a light SSW — is not resolvable with this design, and one proved untestable with the chosen scoring rule; all are reported as findings.
5. **The CFD is not grid-converged at disc altitude** (§10). The magnitude of the sheltering is robust across meshes; the spatial detail of Q within the field is not, and no ranking of launch positions is claimed.

14. Scope and limitations of what was actually computed

This section states the delivered scope plainly, so that no reader has to infer it from the methods.

CFD directions. **Three** inflow directions were computed: $\theta_{\text{wind}} = 90°$, 210° and 270°. The thirteen-direction, 15°-increment sweep described in the project plan (§2) **was not run**. The

three directions were chosen to span the three regimes the study hypothesises — over-lake easterly, prevailing SSW, and over-campus westerly — and they are sufficient to *contrast* those regimes, but they cannot resolve the direction-dependence of Q between them. Any statement in this paper about a "best" or "worst" heading means best or worst *of the three computed*, and is written that way throughout.

Mesh resolution. Production meshes are 4.5×10^5 cells, not the 5–8 $\times 10^6$ of the plan. The refinement is concentrated where the results are read: 2.0×0.50 m cells through the 1–6 m AGL band over the playing surface. Grid independence is demonstrated at that resolution (§10) against a mesh a factor of two coarser in the same region, which bounds discretisation error where it matters; it does not establish that far-field wake structure is grid-converged.

Domain. Upstream, lateral and top extents meet the $5H$ COST 732 guideline; the downstream extent is $10H$ against a $15H$ guideline. Blockage is 0.9%.

Wind speed. A single reference speed, $u_{\mathrm{ref}} = 5\ \mathrm{m\,s^{-1}}$ at 10 m. Because the flow is fully turbulent and the geometry is bluff, the normalised flow field is close to Reynolds-independent, so the Q maps rescale approximately with u_{ref} — but the trajectory model is *not* linear in wind speed, so throwability at 10 $\mathrm{m\,s^{-1}}$ cannot be read off these results.

Validations run and not run. ABL horizontal homogeneity and grid independence were performed and are reported. The **Tominaga (2008) AIJ-B cube benchmark was not run**. **No field measurements were made**, so nothing here is validated against the real site; the anemometer campaign remains future work.

Physical modelling. Steady RANS gives mean-condition playability only: no peak gusts, and no spatially correlated gust structure. The gust model is temporally correlated but spatially uniform along the trajectory at any instant, isotropic (RANS k carries no anisotropy), and Gaussian, so extreme gusts are under-represented. The disc model has no rim side-force term, so predicted crosswind drift is a lower bound. Ground is flat and no tree canopy is included. The hammer throw is exploratory: it spends most of its flight outside the angle-of-attack range over which the tabulated coefficients exist.

15. Author decisions

The open questions carried by the draft are resolved below with defaults that the present results support. Each remains the author's to overrule.

1. **Site.** Sean Earl Field, Loyola Lake Shore Campus, is assumed and used throughout, with the geometry resolved from OSM at 42.001267° N, 87.658214° W. *% FLAG FOR AUTHOR:*

confirm the field outline, which OSM traces as 78 × 25 m — smaller than a regulation Ultimate pitch.

2. **LES.** Out of scope here and recorded as the primary follow-on. The steady-RANS result gives mean-condition playability; it cannot give peak gusts or spatially correlated gust structure, both of which matter for a thrower. An LES of the single easterly case would test precisely the near-wake claim in §11 that RANS is least reliable about. *% FLAG FOR AUTHOR: ∼2 weeks of compute.*

3. **Anemometer validation.** Stated as a future field campaign, not claimed. No field measurements were available for this work, and none are reported. *% FLAG FOR AUTHOR: if the 2026 summer lake-breeze campaign goes ahead, the three field-corner stations proposed in the draft are the right comparison points; the CFD prediction at those points is already in* `flow_fields.npz`*.*

4. **Hammer throw.** Retained but reported as **exploratory**. It is the throw most affected by wind and the one whose coefficients are least well supported: a hammer spends most of its flight well outside the $-10°$ to $+30°$ band over which the tabulated Potts & Crowther coefficients exist, so its trajectories rely on edge-clamped values. Its numbers are reported separately from backhand and forehand and are not used in any headline claim.

5. **Authorship and affiliation.** *% FLAG FOR AUTHOR: author name, department, affiliation, co-authors and funding acknowledgement all require the author's input and are left as placeholders on the title page.*

6. **Venue.** *% FLAG FOR AUTHOR:* on the strength of the present result — a single-site coupled study with a scoped CFD sweep — *Sports Engineering* remains the better fit than splitting into separate trajectory and wind-engineering papers. The disc-model validation (Table 1) is solid enough to stand alone, but the CFD side is currently three directions at one reference wind speed, which is thin for *JWEIA* on its own. Splitting becomes attractive after the full direction sweep and the LES spot-check.

Key references (verified)

Blocken, B. (2014). 50 years of computational wind engineering: Past, present and future. *JWEIA* **129**, 69–102.

Blocken, B. (2018). LES over RANS in building simulation… *Building Simulation* **11**, 821–870.

Crowther, W.J. & Potts, J.R. (2007). Simulation of a spin-stabilised sports disc. *Sports Engineering* **10**(1), 3–21.

Franke, J. et al. (2007). *Best Practice Guideline for CFD Simulation of Flows in the Urban Environment.* COST Action 732.

Hargreaves, D.M. & Wright, N.G. (2007). On the use of the k-ε model in commercial CFD software… *JWEIA* **95**, 355–369.

Hubbard, M. & Hummel, S.A. (2000). Simulation of Frisbee Flight. 5th Math & Computers in Sport, Sydney.

Hummel, S.A. (2003). *Frisbee Flight Simulation and Throw Biomechanics.* MS Thesis, UC Davis.

Kamaruddin, N.M., Potts, J.R. & Crowther, W.J. (2018). Aerodynamic performance of flying discs. *Aircraft Eng. Aerosp. Tech.* **90**(2), 390–397.

Lissaman, P.B.S. & Hubbard, M. (2010). Maximum range of flying discs. *Procedia Engineering* **2**(2), 2529–2535.

Lukes, R., Hart, J. et al. (2022). Disc golf trajectory modelling combining CFD and rigid body dynamics. *Sports Engineering* **25**, 24.

Potts, J.R. & Crowther, W.J. (2002). Frisbee™ aerodynamics. AIAA 2002-3150.

Richards, P.J. & Hoxey, R.P. (1993). Appropriate boundary conditions for computational wind engineering models. *JWEIA* **46–47**, 145–153.

Tominaga, Y. et al. (2008). AIJ guidelines for practical applications of CFD to pedestrian wind environment. *JWEIA* **96**, 1749–1761.

Yang, Y. et al. (2009). New inflow boundary conditions for modelling the neutral equilibrium atmospheric boundary layer. *JWEIA* **97**, 88–95.

Shinozuka, M. & Deodatis, G. (1991). Simulation of stochastic processes by spectral representation. *Applied Mechanics Reviews* **44**(4), 191–204.

Mann, J. (1998). Wind field simulation. *Probabilistic Engineering Mechanics* **13**(4), 269–282.

Garratt, J.R. (1990). The internal boundary layer — a review. *Boundary-Layer Meteorology* **50**, 171–203.

OpenStreetMap contributors (2026). Building footprints, Loyola Lake Shore Campus. https://www.openstreetmap.org (ODbL).

U.S. Geological Survey (2026). 3DEP Elevation Point Query Service. https://epqs.nationalmap.gov.