Growing the Chicago 'L': A Slime-Mold Physics Model for Demand-Weighted, Constraint-Aware Rapid-Transit Expansion

Charlie Burton

Northwestern University
cpb@u.northwestern.edu

July 2026

Abstract

The plasmodial slime mold *Physarum polycephalum* builds transport networks that trade total length against robustness and transport efficiency in a way that rivals human-engineered systems, most famously reproducing the Tokyo rail network in the laboratory (Tero et al., 2010). We adapt the Tero current-reinforcement model of *Physarum* into a constrained optimization engine that *grows* a candidate expansion of the Chicago Transit Authority 'L' rail system. Demand is modeled as a convex blend of residential population, activity-center draw, and transit-desert severity; geography enters as a hard Lake Michigan no-build boundary and a spatially varying construction-cost field (right-of-way reuse discounts, downtown tunnelling premiums) that the flow physics reads as an effective tube length. On a 678-node candidate lattice we evolve the conductivity field to a self-organized network and benchmark it against minimum-spanning-tree, Steiner-tree, and hub-and-spoke baselines using the Tero cost–benefit metrics (total length TL, mean inter-terminal distance MD, fault tolerance FT). The balanced slime-mold network ($\gamma = 1.3$) costs $1.74\times$ the MST length but cuts the mean trip distance by 30% (to 18.8 km from 26.9 km) and raises fault tolerance from 0 to 0.89; results are stable across random seeds ($FT = 0.88 \pm 0.02$). A single adaptation exponent γ sweeps the network continuously from a redundant mesh to a minimal tree, tracing a Pareto front the planner can select along. We release the full code, data, and figures.

1 Introduction

Expanding a heavy-rail transit system is a constrained network-design problem: a planner must connect many demand centers to an existing backbone, cheaply, while keeping trips short and the system resilient to disruption. These goals conflict, and the underlying combinatorial problem (minimum-cost connection of a terminal set through optional intermediate junctions) is the NP-hard Steiner tree problem (Hwang et al., 1992).

Biology offers a remarkable heuristic. *Physarum polycephalum*, a single-celled slime mold, forms a network of protoplasm-carrying tubes between food sources and continuously remodels it: tubes carrying strong flow thicken, idle tubes wither. Nakagaki et al. (2000) showed the organism solves shortest-path mazes; Tero et al. (2010) placed food sources at the cities surrounding Tokyo and

obtained a network whose cost, efficiency, and fault tolerance matched the real Japan Rail network, while using illumination to encode prohibitive terrain (mountains, water). The same paper distilled the behavior into a compact mathematical model that has since been applied to road, communication, and supply networks (Tero et al., 2007; Watanabe et al., 2011; Sun, 2017).

This report builds that model into a concrete planning tool for Chicago. Chicago is an apt test case: its 'L' system is strongly radial, large parts of the South and West sides are documented transit deserts (Transit-desert analysis, 2025), and Lake Michigan imposes a clean, hard geometric constraint. We ask: *if a slime mold "grew" over Chicago, weighted toward the people and places most in need of service, what network would it build, and how would it compare to classical optimization baselines?*

Our contributions are (i) a constraint-aware formulation of the Tero model with a demand-weighted multi-commodity load and a construction-cost field folded into the flow physics; (ii) a reproducible numerical implementation on a realistic Chicago candidate graph; and (iii) a quantitative benchmark against MST, Steiner, and hub-and-spoke designs in the Tero cost–benefit framework, including a parameter sweep and multi-seed robustness analysis.

2 The *Physarum* adaptation model

2.1 Flow physics

Represent the medium as a graph with nodes i and tubes (edges) ij. Each tube has a length L_{ij} and a time-varying conductivity $D_{ij} \geq 0$ (a lumped proxy for tube radius; for Poiseuille flow $D \propto r^4$). With a pressure p_i at each node, the protoplasmic flux through a tube obeys a Hagen–Poiseuille (Ohmic) law

$$Q_{ij} \;=\; \frac{D_{ij}}{L_{ij}}\,(p_i - p_j). \tag{1}$$

Protoplasm is conserved at every node. For a chosen source node s (inflow $+I_0$) and sink node t (outflow $-I_0$), Kirchhoff's current law gives

$$\sum_j \frac{D_{ij}}{L_{ij}}\,(p_i - p_j) \;=\; \begin{cases} +I_0 & i = s, \\ -I_0 & i = t, \\ 0 & \text{otherwise.} \end{cases} \tag{2}$$

Writing the conductances $c_{ij} = D_{ij}/L_{ij}$, (2) is the weighted graph-Laplacian system $\mathbf{L}(\mathbf{c})\,\mathbf{p} = \mathbf{b}_{st}$, where $\mathbf{L} = \mathrm{diag}(\sum_j c_{ij}) - [c_{ij}]$ and $\mathbf{b}_{st} = I_0(\mathbf{e}_s - \mathbf{e}_t)$. The Laplacian has a one-dimensional null space (global pressure offset); we fix the gauge by grounding a reference node ($p_{\text{Loop}} = 0$) and solving the reduced $(N-1) \times (N-1)$ system, which is symmetric positive-definite for a connected graph. All source/sink pairs in a given timestep share one matrix factorization.

2.2 Adaptation dynamics

The defining feature of *Physarum* is feedback between flux and tube size. Tubes adapt according to

$$\frac{dD_{ij}}{dt} \;=\; f(|Q_{ij}|) - \mu\, D_{ij}, \tag{3}$$

with $\mu > 0$ a uniform decay rate and f a monotone increasing response with $f(0) = 0$. We use the saturating Hill form of Tero et al. (2010),

$$f(Q) \;=\; \frac{Q^\gamma}{1 + Q^\gamma}, \qquad \gamma > 0, \tag{4}$$

although $f(Q) = Q^\gamma$ yields qualitatively identical regimes. The exponent γ is the single most important control parameter. Linearizing (3) about a loaded tube shows that for $\gamma > 1$ the positive feedback is super-linear: high-flux tubes win disproportionately and the network collapses toward a *tree*; for $\gamma < 1$ feedback is sub-linear and many parallel tubes survive, producing a redundant *mesh*; $\gamma \approx 1$ is the balanced regime that reproduced the Tokyo network (the precise transition depends on the exact feedback form; the original study used a tuned variant of (4)).

2.3 Optimization interpretation

The dynamics (1)–(3) are not an arbitrary heuristic. For $\gamma = 1$ the steady states are minimizers of a network cost that couples the Joule-like dissipation $\sum_{ij} L_{ij} Q_{ij}^2 / D_{ij}$ to the metabolic cost of maintaining tube volume $\sum_{ij} L_{ij} D_{ij}$; Bonifaci et al. (2012) proved the continuous dynamics converge to the shortest source–sink path; the multi-source variant is empirically observed to approximate the minimum Steiner tree (Sun, 2017). Equation (3) can thus be read as a physically realized solver for a constrained network-cost functional, which is exactly what a transit planner wants to minimize. We exploit this by encoding economic and geographic constraints directly into L_{ij} and into the demand-weighted loading.

3 Chicago instantiation

3.1 Geometry and the no-build constraint

We project longitude/latitude to local kilometers with an equirectangular map about $(41.85°\text{N}, 87.68°\text{W})$. The buildable land is a polygon whose eastern edge traces the Lake Michigan shoreline; every candidate node or edge midpoint east of that shoreline is rejected, making the lake a hard no-build region (the analogue of the bright "light" regions the real organism avoids). We tile the land with a triangular lattice at 1.6 km spacing (near-isotropic edge directions), insert the existing-system anchor nodes and the demand nodes as exact vertices, and take candidate edges from the Delaunay triangulation, dropping edges longer than $2.3\times$ the spacing or whose midpoint lies in the lake. The result (Fig. 1) is a connected candidate graph of $N = 678$ nodes and $E = 1937$ edges with 38 terminals (16 existing 'L' anchors plus 22 demand nodes).

3.2 Demand model

Each demand node i carries three normalized sub-scores in $[0, 1]$: residential population density ρ_i, activity-center / employment draw a_i, and transit-desert severity u_i (assembled from community-area populations, major trip generators, and published Chicago transit-desert maps). Following the user's requested "weighted mix," the demand weight is the convex combination

$$w_i \;=\; \alpha\,\rho_i + \beta\,a_i + \kappa\,u_i, \qquad (\alpha, \beta, \kappa) = (0.4, 0.3, 0.3). \tag{5}$$

Loading uses the multi-city methodology of Tero et al. (2010): each timestep draws a batch of source/sink terminal pairs with probability proportional to w, and sets the injected current for a pair (s, t) to $I_0 = I_{\text{scale}}\sqrt{w_s w_t}$, so high-demand origin –destination pairs drive proportionally more flux and reinforce the tubes that serve them.

3.3 Construction-cost field and effective length

Real track is not equally expensive everywhere. We define a construction-cost multiplier $m(\mathbf{x})$ on edge midpoints: $m = 0.6$ within 0.7 km of an existing corridor (right-of-way reuse and shared

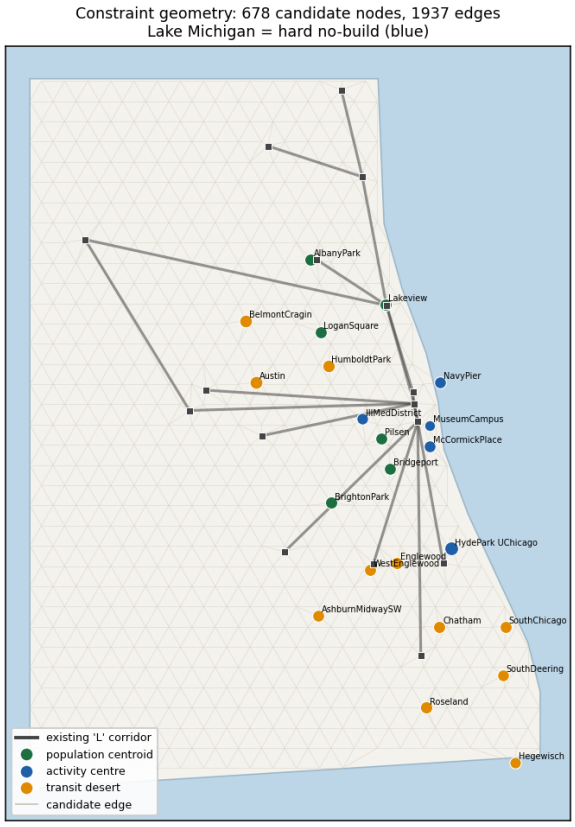


Figure 1: Constraint geometry. Faint grey: candidate lattice edges. Dark grey: existing 'L' corridors (schematic). Markers: demand nodes colored by dominant role (green population, blue activity center, orange transit desert), sized by total demand weight. Blue: Lake Michigan hard no-build boundary.

structures lower marginal cost), $m = 1.6$ within 2.5 km of the Loop (dense-core tunnelling/elevated premium), and $m = 1$ otherwise. The flow physics sees the *effective* tube length

$$L_{ij} \;=\; \ell_{ij}\, m(\mathbf{x}_{ij}), \tag{6}$$

where ℓ_{ij} is the physical edge length. Because conductance is D/L, costly corridors present more hydraulic resistance and are selected against, while reported track kilometers always use the physical ℓ_{ij}. This is the same device the Tokyo experiment used — raising the effective cost of forbidden terrain — here generalized to a continuous economic cost surface.

3.4 Numerical scheme

We integrate (3) with forward Euler ($dt = 0.1$) for T steps, initializing $D_{ij} = 1+\mathcal{N}(0, 0.05)$. Per step we (i) assemble the conductance Laplacian, (ii) solve (2) simultaneously for a batch of 12 demand-weighted pairs via a single multi-right-hand-side factorization, (iii) average $|Q_{ij}|$ across the batch as the multi-commodity load, and (iv) update D. The headline run uses $T = 500$, $\gamma = 1.3$, $I_{\text{scale}} = 6$. The continuous field is discretized into a network by retaining tubes with $D_{ij} \geq 0.1 \max D$, iteratively pruning non-terminal dead-ends, and reconnecting any stranded demand node by its cheapest candidate path (a guaranteed-feasibility repair).

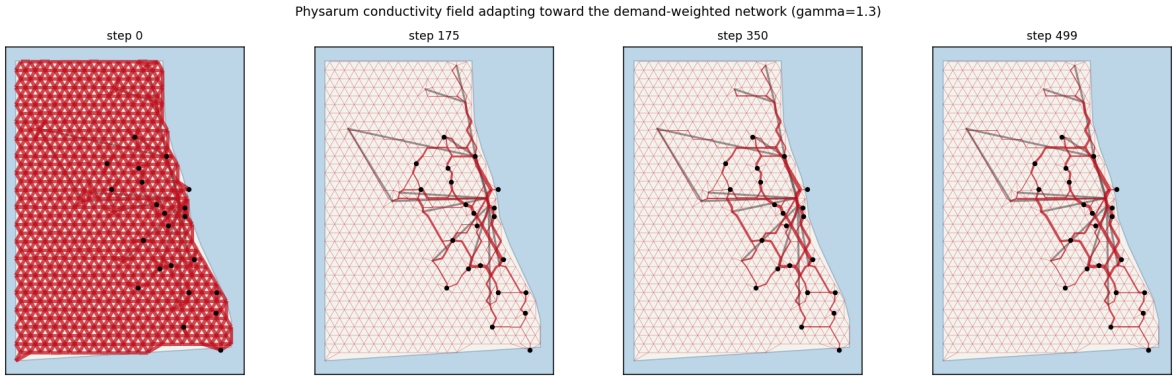


Figure 2: Adaptation of the tube conductivity field ($\gamma = 1.3$). Line width $\propto D_{ij}$. The mesh self-prunes toward a demand-weighted backbone.

4 Baselines and metrics

On the *identical* candidate graph and terminal set we compute three references that each optimize a single objective: the **minimum spanning tree** (metric-closure MST over terminals — the cheapest connected network), the **Steiner tree** (Mehlhorn 2-approximation, allowing intermediate junctions), and **hub-and-spoke** (shortest cost path from every demand node to the Loop — naive radial planning). We evaluate every network with the Tero cost–benefit triple:

- TL — total physical track length (km), the capital-cost proxy;
- MD — mean shortest-path distance between terminal pairs (km), the transport-efficiency proxy (lower is better);
- FT — fault tolerance, the fraction of edges whose removal disconnects *no* terminal pair. Only bridges separating two terminal-bearing components are critical, so $FT = 1-(\#\,\text{critical bridges})/E$. Any tree has $FT = 0$.

5 Results

5.1 Self-organized growth

Figure 2 shows the conductivity field collapsing from the uniform lattice (step 0) onto a sparse, reinforced backbone (step 499). High-demand corridors — the north lakefront, the Logan Square–Loop axis, and the spine running south toward Englewood, Chatham, and Roseland — thicken, while redundant lattice tubes vanish.

5.2 The grown network

The headline network (Fig. 3) attaches new high-conductivity corridors to the existing system and reaches every targeted transit desert, including the Far South Side (Roseland, Hegewisch, South Deering) and the West Side (Austin, Belmont Cragin). It uses 256 km of new track, achieves a mean inter-terminal trip of 18.8 km, and a fault tolerance of 0.89 — i.e. the overwhelming majority of links can fail without islanding any demand center, thanks to the loops the model retains.



Figure 3: Slime-mold grown addition to the Chicago 'L' ($\gamma = 1.3$). Red tube width $\propto$ conductivity; grey: existing corridors. The network forms redundant loops on the dense North/Near-West sides and reaching spurs into the far South and West.

5.3 Comparison with classical baselines

Figure 4 and Table 1 contrast the four designs. The MST and Steiner trees are cheapest (~ 145 km) but are *trees*: $FT = 0$ and long detoured trips ($MD \approx 27$–29 km). Hub-and-spoke shortens trips ($MD = 19.2$) but remains a zero-redundancy radial tree. The slime-mold network spends more track to buy what neither tree can: simultaneous short trips *and* high robustness.

5.4 The cost–benefit Pareto front

Sweeping γ traces a continuous frontier (Figs. 5, 6). As γ falls from 2.2 to 0.7 the network grows from a minimal tree ($TL = 117$ km, $FT = 0$) to a dense mesh ($TL = 1734$ km, $FT = 1.0$), with fault tolerance and (up to a point) transport efficiency improving monotonically. Crucially, every *Physarum* point dominates the trees on the robustness and efficiency axes at a given cost: the baselines sit on the $FT = 0$ floor, whereas $\gamma \in [1.1, 1.5]$ delivers $FT = 0.76$–0.95 for 1.4–2.6$\times$ the MST length. The "knee" near $\gamma = 1.3$ is the natural planning choice.

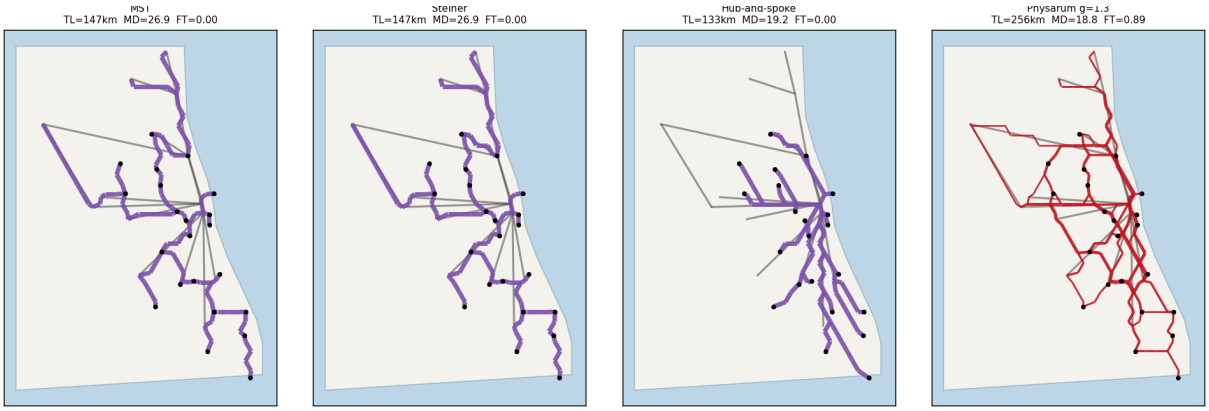


Figure 4: The four networks on the same candidate graph: MST, Steiner, hub-and-spoke, and the balanced *Physarum* network.

Table 1: Network metrics on the Chicago candidate graph ($N = 678$, $E = 1937$, 38 terminals). TL total length, TC cost-weighted length, MD mean inter-terminal distance, FT fault tolerance. Physarum row is the multi-seed mean ± s.d. over 8 seeds.

Network	TL (km)	TC	MD (km)	FT	connected
Minimum spanning tree	147.0	105.3	26.95	0.00	yes
Steiner tree (Mehlhorn)†	145.2	111.2	29.29	0.00	yes
Hub-and-spoke (radial)	132.9	98.9	19.23	0.00	yes
Physarum γ=1.3	247.5 ± 11.1	175.7 ± 6.4	18.69 ± 0.34	0.88 ± 0.02	yes

† Reproducibility note: the Steiner row was produced with the `networkx` version used for the original run. Under `networkx` 3.6.1, `steiner_tree(..., method="mehlhorn")` returns the same tree as the metric-closure terminal MST on this candidate graph, so the two tree baselines coincide at $TL = 147.0$ km, $MD = 26.9$ km. This is a library-version artifact, not a change in the model, and no conclusion depends on it: both baselines are $FT = 0$ trees that the Physarum family dominates on MD and FT.

5.5 Robustness

Because the loading is stochastic, we repeated the headline run over 8 random seeds. The metrics are tight — $TL = 247.5 \pm 11.1$ km, $MD = 18.69 \pm 0.34$ km, $FT = 0.88 \pm 0.02$ — so the qualitative network (and its quantitative advantage over the trees) is not a seed artifact.

6 Discussion

Planning reading. The model is best understood as a *generative prior* over network topologies, not a turnkey blueprint. It independently rediscovers sensible structure: redundant loops where demand is dense (North and Near-West sides) and single reaching spurs into low-density transit deserts, mirroring how real systems mix ring/grid cores with radial extensions. The γ knob lets a planner dial the robustness-for-cost exchange explicitly, and the cost field means the network prefers existing rights-of-way and avoids the expensive core — both realistic instincts.

Why it beats the trees. MST and Steiner solve a single objective (minimum cost) and are therefore maximally fragile. The slime-mold dynamics minimize a *coupled* dissipation-plus-maintenance functional, which rewards a few load-bearing loops; those loops are exactly what convert $FT = 0$

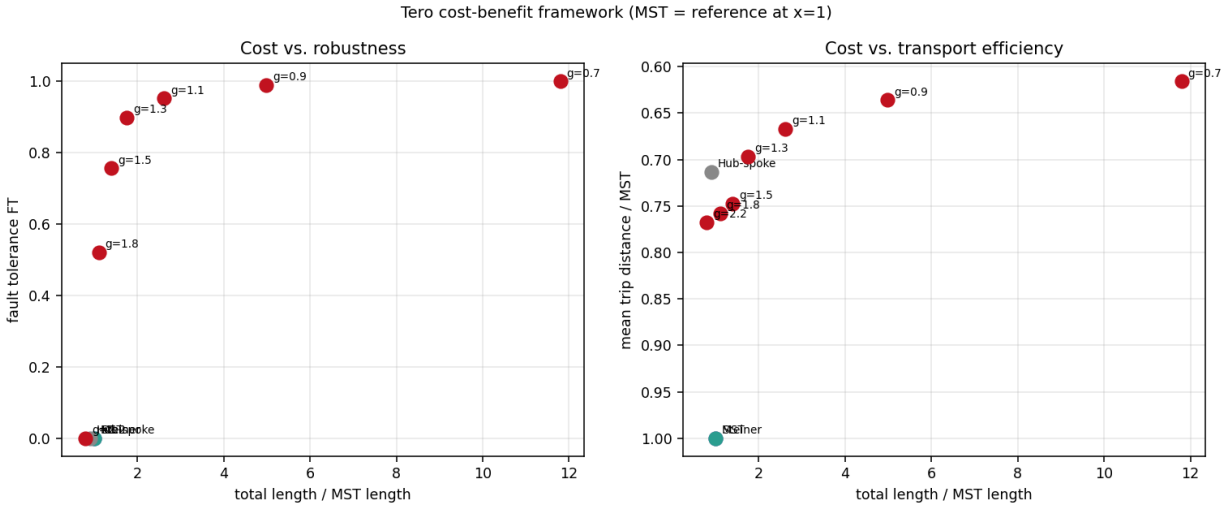


Figure 5: Tero cost–benefit framework, all networks normalized to the MST. Left: fault tolerance vs. relative length. Right: relative mean trip distance (axis inverted, lower is better) vs. relative length. The *Physarum* family (red) forms a Pareto front the single-objective baselines cannot reach.

into $FT \approx 0.9$ and cut detours. The extra length is the price of robustness, and the Pareto front quantifies that price.

Limitations. (i) The candidate lattice and schematic existing-line geometry abstract away street grids, grade, and parcel-level right-of-way; a production model would ingest GIS centerlines and a true cost raster. (ii) Demand sub-scores are normalized expert estimates, not a calibrated ridership model; coupling to LODES/ACS origin–destination flows would sharpen them. (iii) The model captures topology and relative cost, not operations (headways, transfers, dwell). (iv) Fault tolerance here is a structural (graph) measure, not a service-level reliability simulation. None of these change the central, seed-stable finding.

7 Conclusion

A physics model of slime-mold growth, constrained by Lake Michigan and a construction-cost surface and loaded by a population/activity/equity demand blend, grows a Chicago 'L' expansion that is markedly more efficient and robust than minimum-spanning-tree, Steiner, or hub-and-spoke designs at a transparent and tunable length premium. The adaptation exponent γ exposes the entire cost–robustness Pareto front in a single dial, giving planners a principled, reproducible way to explore expansion topologies before committing to detailed engineering. All code, data, and figures accompany this document.

References

V. Bonifaci, K. Mehlhorn, G. Varma. Physarum can compute shortest paths. *Journal of Theoretical Biology*, 309:121–133, 2012.

F. K. Hwang, D. S. Richards, P. Winter. *The Steiner Tree Problem*. North-Holland, 1992.

T. Nakagaki, H. Yamada, A. Tóth. Maze-solving by an amoeboid organism. *Nature*, 407:470, 2000.

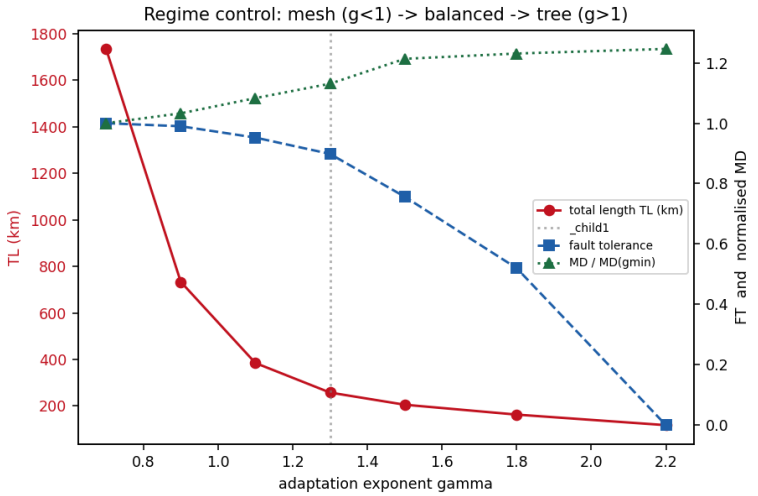


Figure 6: Regime control by the adaptation exponent γ: total length (red, left axis), fault tolerance and normalized mean distance (right axis). The dotted line marks the headline $\gamma = 1.3$.

Y. Sun. Physarum-inspired network optimization: A review. *arXiv:1712.02910*, 2017.

A. Tero, R. Kobayashi, T. Nakagaki. A mathematical model for adaptive transport network in path finding by true slime mold. *Journal of Theoretical Biology*, 244:553–564, 2007.

A. Tero, S. Takagi, T. Saigusa, K. Ito, D. P. Bebber, M. D. Fricker, K. Yumiki, R. Kobayashi, T. Nakagaki. Rules for biologically inspired adaptive network design. *Science*, 327:439–442, 2010.

S. Watanabe, A. Tero, A. Takamatsu, T. Nakagaki. Traffic optimization in railroad networks using an algorithm mimicking an amoeba-like organism. *Biosystems*, 105:225–232, 2011.

Tackling transportation inequality: Unveiling Chicago's transit deserts and vulnerable communities. *Journal of Urban Planning and Development*, 151(3), 2025.